

HCLSoftware

Software Containerization Lesson 2


l.ziosi@vu.nl, lara.ziosi@hcl.com

Agenda

- Kubernetes architecture
 - Control plane (Master nodes)
 - Worker nodes
 - kubectl command
- Nodes
- Pods and containers
- Namespaces
- Creating pods with kubectl run
- YAML syntax
- Deployment
- ReplicaSet
- DaemonSet
- Job
- CronJob
- Garbage Collector

Kubernetes Architecture – control plane

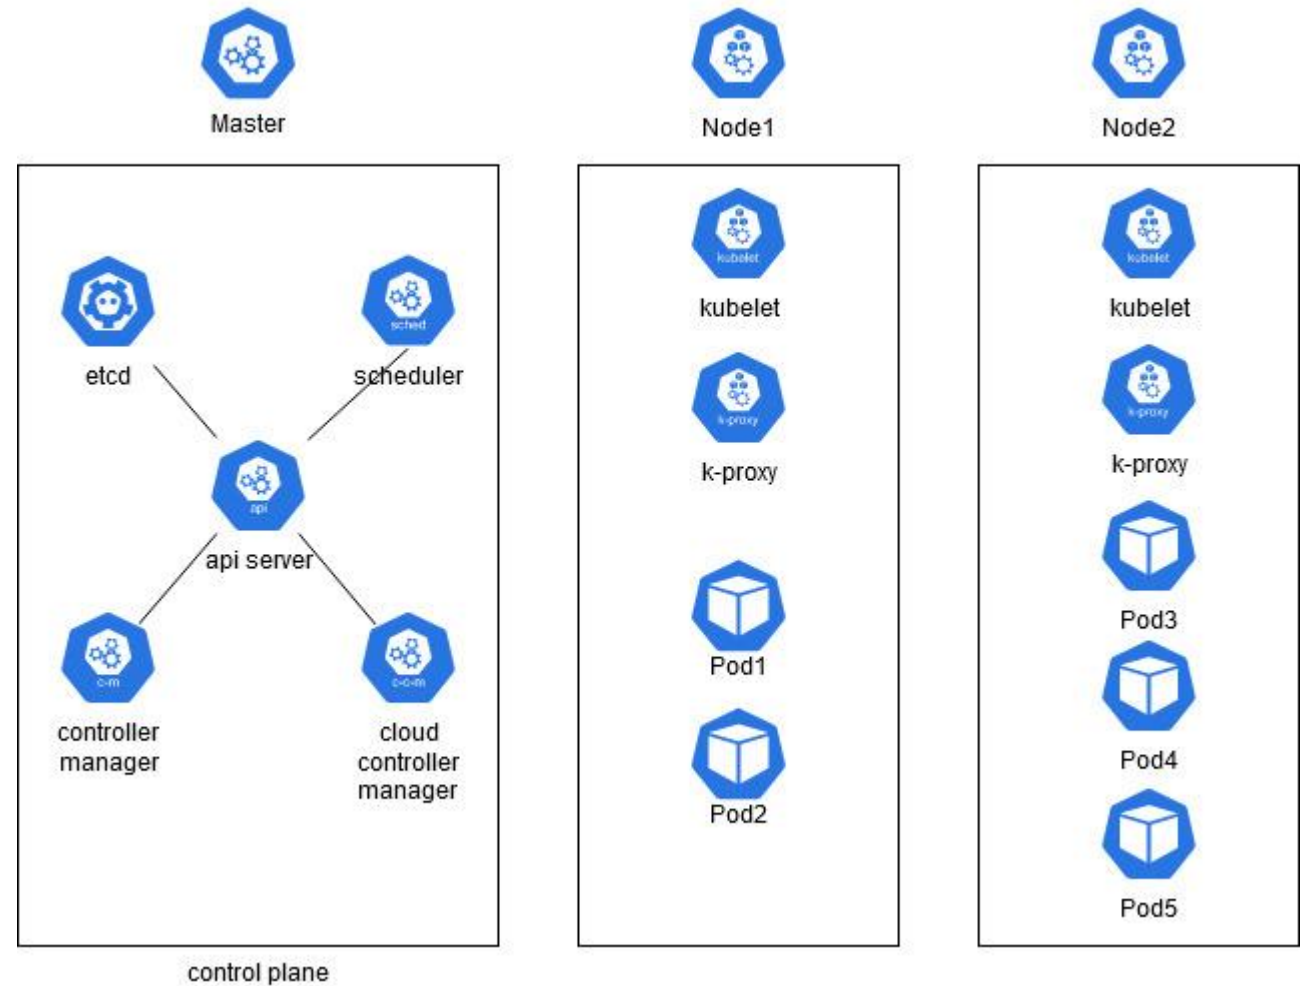
- **Control plane:** 1-3+ Master nodes
- **Master Node:** Physical or Virtual Machine that doesn't run user workloads directly but manages the cluster
- **kube-apiserver:** expose the Kubernetes api
- **etcd:** storage of key value pairs for Kubernetes configuration
- **kube-scheduler:** selects the node on which to run newly created pods
- **kube-controller-manager:** runs controller processes (Node, Replication, Endpoints, Service Account & Token controllers)
- **Cloud controller manager:** connects the cluster to the cloud provider's API, separating components that interact with the cloud platform from components that interact with the kubernetes cluster.



Reference: <https://kubernetes.io/docs/concepts/overview/components/>

Kubernetes Architecture – Worker nodes

- **Node/Worker/Minion:** Physical or Virtual Machine that executes user workloads, interacting with the control plane
- **kubelet:** ensures that pods are running in the node as declared by their specifications
- **kube-proxy:** network proxy, maintains network rules on nodes
- **pod:** group of one or more containers that share storage/network resources, described by a specification about how to run the containers



<https://kubernetes.io/docs/concepts/overview/components/>

Kubectl command

- The kubectl command lets you interact with the kubernetes cluster. It is a utility that invokes the kubernetes api.
- The general syntax of kubectl is:

```
kubectl [command] [TYPE] [NAME] [flags]
```

- The value of **command** can be: create, get, describe, delete, logs, exec, run and many more
- The value of **type** can be any resource type (examples: node, pod, service, namespace, configmap, secret, job). Types are case-insensitive and you can specify the singular, plural, or abbreviated forms
- The value of **name** is case sensitive and if it is omitted, then the command will apply to all objects of that type
- The value of **flags** can be obtained using `kubectl help`
- Examples:

```
kubectl get pods
kubectl delete pod <pod-name>
kubectl describe pod <pod-name>
kubectl get pods --namespace <namespace-name>
```

- You can abbreviate types, for example write pod instead of pods: `kubectl get pod`

<https://kubernetes.io/docs/reference/kubectl/overview/>

Kubectl configuration

Kubectl needs to be configured so it can connect to the control-plane. The configuration is stored in the file:

```
~/.kube/config
```

If you use a microk8s single node, you don't need to set up the configuration manually.

You can view the configuration by running the command:

```
microk8s config
```

For more information, see:

<https://microk8s.io/docs/working-with-kubectl>

```
lara@kube-master-gui:~/k8s/jobs$ microk8s config
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tck1JSURBVENDQW
  4TkRBNU16VmFGdzB6TURFeU1qa3hOREE1TXpwYQpNQmN4RLRBVEJnTLZCQU1NRERFd0xqRTFNATR4T0RNd
  JVbmxNY3AKTDR5aVlwktSWDgycUd3d1BaeEVCaDh0TEJLd3Y0TUhlbFdpRTZ5bFI2TmxWZELWUmFLYlN5
  XFyeXNBNUFNHdIUvLEVLIWt0JCWUUVGTzZwClhzN0ZQVTFydTNQcWFrNWh4U2dDODLSd01C0EdBMVvkSX
  d0VBQWFOUU1FNHdIUvLEVLIWt0JCWUUVGTzZwClhzN0ZQVTFydTNQcWFrNWh4U2dDODLSd01C0EdBMVvkSX
  DUWg5LWpNZnpFdkVEa2tWOFU3Z3ZubUNqbzIrWjFqMGY2b3lkVku4TnVEdVJOVDNQemtKNnY0aDRkeEJCm
  swTnAvTEVEQ3VOUULBM3EzWDUwZThHc3QxUnFSMHJvN2d4a1V6OXVvSEdQZEkzTgpYamNCV1FNZ0pmWELL
  VptMTN0bEhWWldPT0prZHBUR2h5SElRMFGkbHluVVIQT0KLS0tLS1FTkQgQ0VSVElGSUNBVEUtLS0tLQo
  server: https://10.0.2.15:16443
  name: microk8s-cluster
contexts:
- context:
  cluster: microk8s-cluster
  user: admin
  name: microk8s
current-context: microk8s
kind: Config
preferences: {}
users:
- name: admin
  user:
    token: M2JJMk11NHNWcUtPTFpMNzNWaHh0K09vb1locWppYi9uMzFGS0ZneUhRaz0K
```


Kubectl examples

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
minikube	Ready	master	6m23s	v1.17.3

```
$ kubectl get namespace
```

NAME	STATUS	AGE
default	Active	6m57s
kube-node-lease	Active	6m59s
kube-public	Active	6m59s
kube-system	Active	6m59s
kubernetes-dashboard	Active	6m50s

```
$ kubectl get pods --namespace kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-6955765f44-wr4bv	1/1	Running	0	35s
coredns-6955765f44-z8qwl	1/1	Running	0	35s
etcd-minikube	1/1	Running	0	38s
kube-apiserver-minikube	1/1	Running	0	38s
kube-controller-manager-minikube	1/1	Running	0	38s
kube-proxy-tp2cx	1/1	Running	0	35s
kube-scheduler-minikube	1/1	Running	0	38s
storage-provisioner	1/1	Running	0	39s

```
$
```

Powered by  KataCoda

```
$ kubectl describe pod etcd-minikube
```

```
Error from server (NotFound): pods "etcd-minikube" not found
```

```
$ kubectl describe pod etcd-minikube -n kube-system
```

```
Name:          etcd-minikube
Namespace:     kube-system
Priority:       2000000000
Priority Class Name: system-cluster-critical
Node:          minikube/172.17.0.22
Start Time:    Sat, 14 Nov 2020 13:40:01 +0000
```

Powered by  KataCoda

```
$ kubectl get services -n kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP,9153/TCP	3m53s

```
$ kubectl get secrets -n kube-system
```

NAME	TYPE
attachdetach-controller-token-mrmhv	kubernetes.io/service-account-token
bootstrap-signer-token-fknsd	kubernetes.io/service-account-token
bootstrap-token-a53kyo	bootstrap.kubernetes.io/token
certificate-controller-token-pdsv5	kubernetes.io/service-account-token
clusterrole-aggregation-controller-token-9sk2k	kubernetes.io/service-account-token
coredns-token-5w8zd	kubernetes.io/service-account-token
cronjob-controller-token-rkk78	kubernetes.io/service-account-token
daemon-set-controller-token-8xjf7	kubernetes.io/service-account-token
default-token-ckc4h	kubernetes.io/service-account-token
deployment-controller-token-85n6v	kubernetes.io/service-account-token
disruption-controller-token-482n8	kubernetes.io/service-account-token
endpoint-controller-token-s4kqj	kubernetes.io/service-account-token

Try it yourself in the browser at: <https://kubernetes.io/docs/tutorials/kubernetes-basics/create-cluster/>

Nodes

- A Node is a physical or virtual machine that runs the **kubelet** agent, the **kube-proxy** pod and **Network pod**.
- **Kubelet installation:**
 - If you install k8s with **kubeadm**, kubelet must be installed manually on each worker node as described in:
<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>
 - If you install k8s using **microk8s**, kubelet will be installed automatically.
- **Network pod:** Kubernetes does not include any Network pod. You need to choose and install a network add-on.
 - In the following example, the Network pod is provided by **Weave Net** but there are other options:
<https://www.weave.works/docs/net/latest/kubernetes/kube-addon/>
 - If you install **microk8s**, the **Calico** Network is installed by default.
- These components allow the cluster to manage the pods that run on the node.
- Each worker Node needs to be registered with the cluster.

Node example with WeaveNet Network

```
[root@kbnodel ~]# kubectl describe node kbnodel
Name: kbnodel
Roles: <none>
Labels: beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/os=linux
        kubernetes.io/arch=amd64
        kubernetes.io/hostname=kbnodel
        kubernetes.io/os=linux
Annotations: kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock
              node.alpha.kubernetes.io/ttl: 0
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Fri, 04 Sep 2020 19:23:04 +0530
Taints: <none>
Unschedulable: false
Lease:
  HolderIdentity: kbnodel
  AcquireTime: <unset>
  RenewTime: Sun, 15 Nov 2020 15:31:49 +0530
Conditions:
  Type           Status  LastHeartbeatTime               LastTransitionTime             Reason                       Message
  ----           -
  NetworkUnavailable  False   Sun, 15 Nov 2020 14:46:25 +0530   Sun, 15 Nov 2020 14:46:25 +0530   WeaveIsUp                   Weave pod has set this
  MemoryPressure      False   Sun, 15 Nov 2020 15:28:05 +0530   Fri, 04 Sep 2020 19:23:05 +0530   KubeletHasSufficientMemory   kubelet has sufficient memory available
  DiskPressure        False   Sun, 15 Nov 2020 15:28:05 +0530   Fri, 04 Sep 2020 19:23:05 +0530   KubeletHasNoDiskPressure     kubelet has no disk pressure
  PIDPressure         False   Sun, 15 Nov 2020 15:28:05 +0530   Fri, 04 Sep 2020 19:23:05 +0530   KubeletHasSufficientPID      kubelet has sufficient PID available
  Ready              True    Sun, 15 Nov 2020 15:28:05 +0530   Fri, 04 Sep 2020 19:24:35 +0530   KubeletReady                 kubelet is posting ready status
Addresses:
  InternalIP: 10.0.2.7
  Hostname: kbnodel
Capacity:
  cpu: 2
  ephemeral-storage: 45615348Ki
  hugepages-2Mi: 0
  memory: 8000364Ki
  pods: 110
Allocatable:
  cpu: 2
  ephemeral-storage: 42039104648
  hugepages-2Mi: 0
  memory: 7897964Ki
  pods: 110
```

Node example continued

```
System Info:
Machine ID:          9cfc10c654e9413e8116328ba497b2a3
System UUID:         61b81414-9ad9-400c-b804-5c83c6a90209
Boot ID:             c54b6a8c-8449-407e-bc28-5507f3683d98
Kernel Version:      4.18.0-147.el8.x86_64
OS Image:            CentOS Linux 8 (Core)
Operating System:    linux
Architecture:        amd64
Container Runtime Version: docker://18.9.1
Kubelet Version:      v1.19.0
Kube-Proxy Version:  v1.19.0
```

```
Non-terminated Pods: (19 in total)
```

Namespace	Name	CPU Requests	CPU Limits	Memory Requests	Memory Limits	AGE
default	details-v1-79c697d759-vsxhw	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
default	httpbin-74fb669cc6-4npjs	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	33d
default	nginx-app-d6ff45774-grm5d	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	42m
default	productpage-v1-65576bb7bf-zrk6g	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
default	ratings-v1-7d99676f7f-v795f	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
default	reviews-v1-987d495c-5jqxd	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
default	reviews-v2-6c5bf657cf-49zfp	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
default	reviews-v3-5f7b9f4f77-w9x4h	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
istio-system	grafana-57bb676c4c-zxpqr	0 (0%)	0 (0%)	0 (0%)	0 (0%)	37d
istio-system	istio-egressgateway-b9d46896-qf572	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
istio-system	istio-ingressgateway-dc76747bf-vqqrg	10m (0%)	2 (100%)	40Mi (0%)	1Gi (13%)	71d
istio-system	istiod-69c88fcb8-kprb6	10m (0%)	0 (0%)	100Mi (1%)	0 (0%)	71d
istio-system	jaeger-75948789b4-k4qh6	10m (0%)	0 (0%)	0 (0%)	0 (0%)	37d
istio-system	kiali-7d5cb68b45-ptd44	0 (0%)	0 (0%)	0 (0%)	0 (0%)	37d
istio-system	prometheus-7c8bf6df84-l7bpj	0 (0%)	0 (0%)	0 (0%)	0 (0%)	37d
kube-system	kube-proxy-m7glv	0 (0%)	0 (0%)	0 (0%)	0 (0%)	71d
kube-system	weave-net-jm29z	100m (5%)	0 (0%)	200Mi (2%)	0 (0%)	71d
sftp	sftp-768748c469-qflfr	0 (0%)	0 (0%)	0 (0%)	0 (0%)	32d
sftp	sftp-66b57db894-nqspn	0 (0%)	0 (0%)	0 (0%)	0 (0%)	33d

Node example continued

```
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource           Requests      Limits
-----
cpu                 220m (11%)    20 (1000%)
memory              700Mi (9%)    10Gi (132%)
ephemeral-storage   0 (0%)        0 (0%)
hugepages-2Mi       0 (0%)        0 (0%)

Events:
Type      Reason                Age          From          Message
-----
Normal    Starting              46m          kubelet, kbnode1    Starting kubelet.
Normal    NodeAllocatableEnforced 46m          kubelet, kbnode1    Updated Node Allocatable limit across pods
Normal    NodeHasNoDiskPressure  46m (x7 over 46m) kubelet, kbnode1    Node kbnode1 status is now: NodeHasNoDiskPressure
Normal    NodeHasSufficientPID   46m (x7 over 46m) kubelet, kbnode1    Node kbnode1 status is now: NodeHasSufficientPID
Warning   Rebooted              46m          kubelet, kbnode1    Node kbnode1 has been rebooted, boot id: c54b6a8c-8449-407e-bc28-5507f3683d98
Normal    NodeHasSufficientMemory 46m (x8 over 46m) kubelet, kbnode1    Node kbnode1 status is now: NodeHasSufficientMemory
Normal    Starting              46m          kube-proxy, kbnode1 Starting kube-proxy.
```

```
[root@kbnode1 ~]# ps -ef |grep kubelet
root      1498      1  9 14:44 ?        00:05:01 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml --network-plugin=cni --pod-infra-container-image=k8s.gcr.io/pause:3.2
```

If you install microk8s, the installation paths are different:
/snap/microk8s/1864/kubelet
/snap/microk8s/1864/kube-apiserver

Communication between control plane and worker nodes

If you work in a multi-node cluster where restrictive firewall policies are enforced on each node, you need to modify the firewall rules and open all the ports required for the communication needs of the control planes and of the worker nodes.

Control-plane node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	6443*	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10251	kube-scheduler	Self
TCP	Inbound	10252	kube-controller-manager	Self

Worker node(s) [↔](#)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	30000-32767	NodePort Services†	All

† Default port range for [NodePort Services](#).

Table from: <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

Namespaces

- Namespaces are used to create virtual clusters inside a cluster to avoid name conflicts, restrict resource usage establish quotas, attach policies that define who can perform which actions.
- By default, K8s creates objects in the **default** namespace.
- K8s components run in separate namespace called **kube-system**.
- You can create additional namespaces to separate the work of different teams or different applications:

```
kubectl create namespace <namespace-name>
```

- Namespaces cannot be nested.
- Not all resource types belong to a namespace. For example, pods and services are in a namespace. Namespaces, persistent volumes and nodes aren't. To see resource types that are (or aren't) in a namespace use:

```
kubectl api-resources --namespaced=true  
kubectl api-resources --namespaced=false
```

- To see which resources of a certain type are in a namespace, append `-n namespace-name`:

```
kubectl get pods -n <namespace-name>  
kubectl describe pod <pod-name> -n <namespace-name>
```

Namespace examples

The following outputs are truncated:

```
[root@kbnodel ~]# kubectl create namespace sample
namespace/sample created
```

```
[root@kbnodel ~]# kubectl describe namespace sample
Name:         sample
Labels:       <none>
Annotations:  <none>
Status:       Active

No resource quota.

No LimitRange resource.
```

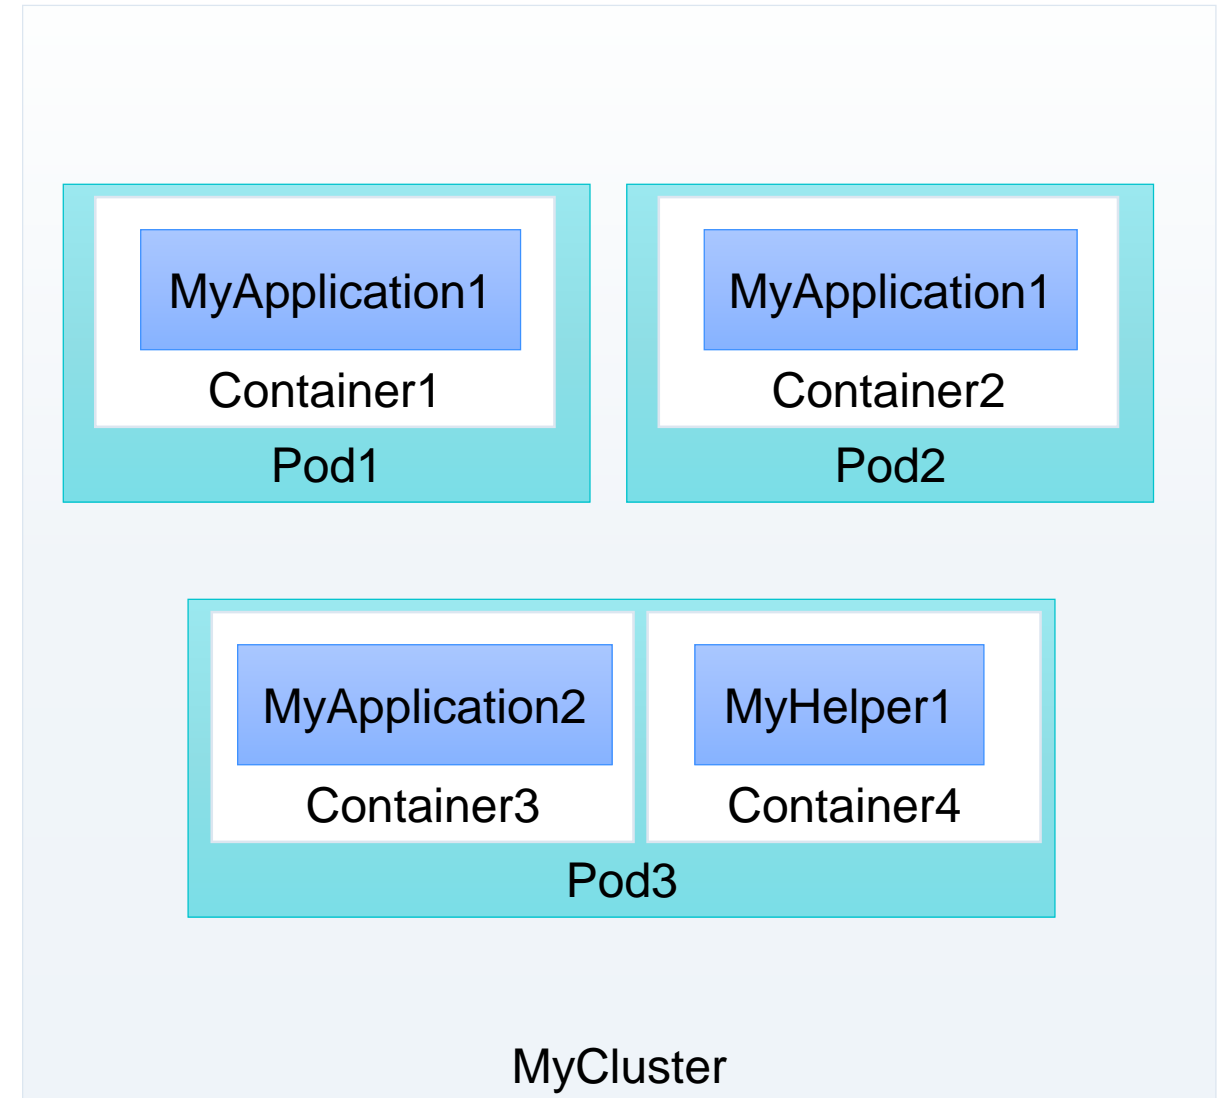
```
[root@kbnodel ~]# kubectl get pods -n sample
No resources found in sample namespace.
[root@kbnodel ~]# kubectl get pods -n default
NAME                                READY   STATUS    RESTARTS   AGE
details-v1-79c697d759-vsxhw         2/2     Running   8           71d
httpbin-74fb669cc6-4npjs            2/2     Running   4           33d
nginx-app-d6ff45774-grm5d           2/2     Running   0           33m
productpage-v1-65576bb7bf-zrk6g     2/2     Running   8           71d
ratings-v1-7d99676f7f-v795f         2/2     Running   8           71d
reviews-v1-987d495c-5jqxd            2/2     Running   8           71d
reviews-v2-6c5bf657cf-49zfp         2/2     Running   8           71d
reviews-v3-5f7b9f4f77-w9x4h         2/2     Running   8           71d
```

```
[root@kbnodel ~]# kubectl api-resources --namespaced=true
NAME                SHORTNAMES  APIGROUP  NAMESPACED  KIND
bindings            cm          apps      true        Binding
configmaps           cm          core      true        ConfigMap
endpoints            ep          core      true        Endpoints
events               ev          core      true        Event
limitranges          limits      core      true        LimitRange
persistentvolumeclaims pvc         core      true        PersistentVolumeClaim
pods                 po          core      true        Pod
podtemplates         po          core      true        PodTemplate
replicationcontrollers rc           core      true        ReplicationController
resourcequotas       quota       core      true        ResourceQuota
secrets              sa          core      true        Secret
serviceaccounts      sa          core      true        ServiceAccount
services             svc         core      true        Service
controllerrevisions  cr          apps      true        ControllerRevision
daemonsets           ds          apps      true        DaemonSet
deployments          deploy      apps      true        Deployment
replicasets          rs          apps      true        ReplicaSet
statefulsets         sts         apps      true        StatefulSet
```

```
[root@kbnodel ~]# kubectl api-resources --namespaced=false
NAME                SHORTNAMES  APIGROUP  NAMESPACED  KIND
componentstatuses   cs          core      false       ComponentStatus
namespaces           ns          core      false       Namespace
nodes               no          core      false       Node
persistentvolumes   pv          core      false       PersistentVolume
```

Pods and Containers

- A **pod** is the smallest unit you can deploy to a Kubernetes Cluster
- A pod generally contains a **single container**. To **scale horizontally**, you create multiple identical pods.
- A pod may contain **multiple containers** that are strictly related. All containers in a pod share the same network and storage resources.
- A pod may contain one or more **init containers** used to initialize the application. These containers run in sequence and end as soon as they finish their job.
- You can create pods directly from the command line but more often you let the **kube-controller** create them for you.
- In this case you provide the pod specification as part of a **Deployment, ReplicaSet, StatefulSet, DaemonSet or Job**.



Creating a pod from the command line

To create and run a pod from the command line you use the `kubectl run` command.

Example:

```
kubectl run -image <image-name> <pod-name> --port=<port_on_container_ip> --hostport=<port_on_host_ip>
```

```
lara@kube-master-gui:~$ kubectl run --image nginx nginx-app --port=80 --hostport=80
pod/nginx-app created
```

To check that the pod is running and see its IP address you can use:

```
kubectl get pods -o wide
```

```
lara@kube-master-gui:~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
nginx-app	1/1	Running	0	30s	10.1.96.187	kube-master-gui	<none>		<none>	

For more options of the `kubectl run` command, type:

```
kubectl run --help
```

Pod lifecycle

The lifecycle of a pod goes through different phases, that can take the following values:

Order	Phase	Description
1	Pending	Cluster accepted the Pod, but not all containers are running yet. Possible reasons: scheduling is still ongoing or missing images are being downloaded
2	Running	The pod is bound to a node. All containers have been created. At least one container is running, starting or restarting
3.1	Succeeded	All containers terminated successfully; they will not restart.
3.2	Failed	All containers terminated. At least one container terminated with error.
any	Unknown	Status cannot be detected (controller cannot communicate with worker node)

If the status is **CrashLoopBackoff**, there's a major problem with the configuration of the pod and one of its containers exited unexpectedly.

Getting the details of the running pod

- To get the IP address of the running pod, you can use:
`kubectl describe pod <pod-name>`
- In the example below you can see that the IP address is: 10.1.96.170. If the pod stops and restarts, the IP will change.
- You can also see the Port and Host Port values associated with a container.
- The Container State can be: **Waiting**, **Running** or **Terminating**.
- If the container image fails to download the State will be **Waiting** with reason: `ImagePullBackoff`.

```
lara@kube-master-gui:~$ kubectl describe pod nginx-app
Name:          nginx-app
Namespace:     default
Priority:       0
Node:          kube-master-gui/10.0.2.15
Start Time:    Sun, 03 Jan 2021 15:16:57 +0100
Labels:        run=nginx-app
Annotations:   cni.projectcalico.org/podIP: 10.1.96.170/32
               cni.projectcalico.org/podIPs: 10.1.96.170/32
Status:        Running
IP:            10.1.96.170
IPs:
  IP: 10.1.96.170
Containers:
  nginx-app:
    Container ID:  containerd://33eaf6dd17c0e7910b9ceed049cdbacdccaf91081c864260729f7caa1a4a0c22
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:4cf620a5c81390ee209398ecc18e5fb9dd0f5155cd82adcbae532fec94006fb9
    Port:          80/TCP
    Host Port:     80/TCP
    State:         Running
      Started:     Sun, 03 Jan 2021 15:17:01 +0100
    Ready:         True
    Restart Count:  0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-tnd79 (ro)
```

Checking that the nginx service is running

If your pod has exposed a port, you can check for the nginx service using the current IP address of the pod and the port:

```
lara@kube-master-gui:~$ curl http://10.1.96.176:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

If your pod has exposed a hostport, you can check for the nginx service using the IP address of the host and the hostport:

```
lara@kube-master-gui:~$ curl http://localhost:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

Running a command inside a container

To run a shell inside a container (that has that shell installed), you can use a command like:

```
kubectl exec -it <pod-name> -- /bin/bash
```

You might need to use a different shell, for example use `ash` in Alpine Linux.

The `-t` flag tells the container that the input is from a TTY and `-i` ensures that the shell input is passed to the container. To exit from shell, type `exit`.

You can also run other commands that are available inside the container.

In the example below, you can see where the `nginx` image configures the server to listen to port 80.

```
lara@kube-master-gui:~$ kubectl exec -it nginx-app -- /bin/bash
root@nginx-app:/# ls
bin    docker-entrypoint.d  home  media  proc  sbin  tmp
boot  docker-entrypoint.sh lib    mnt    root  srv   usr
dev    etc                  lib64  opt    run   sys   var
root@nginx-app:/# more /etc/nginx/conf.d/default.conf
server {
    listen      80;
    listen  [::]:80;
    server_name localhost;
```

Getting the logs of a pod

While the pod runs, you can view the logs using the command:

```
kubectl logs <pod_name> -n <namespace_name> -c <container_name>
```

The log combines the stdout and stderr of the applications running in the pod.

To specify a container inside the pod, use the `-c` flag followed by the container name.

```
lara@kube-master-gui:~$ kubectl logs nginx-app
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
10.0.2.15 - - [03/Jan/2021:20:09:24 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.68.0" "-"
lara@kube-master-gui:~$
```

HCLSoftware

Kubernetes Workloads

Workload resources and controllers

Typically, instead of running individual pods, you create *workload resources* that manage a set of pods based on the *desired state* that you declare.

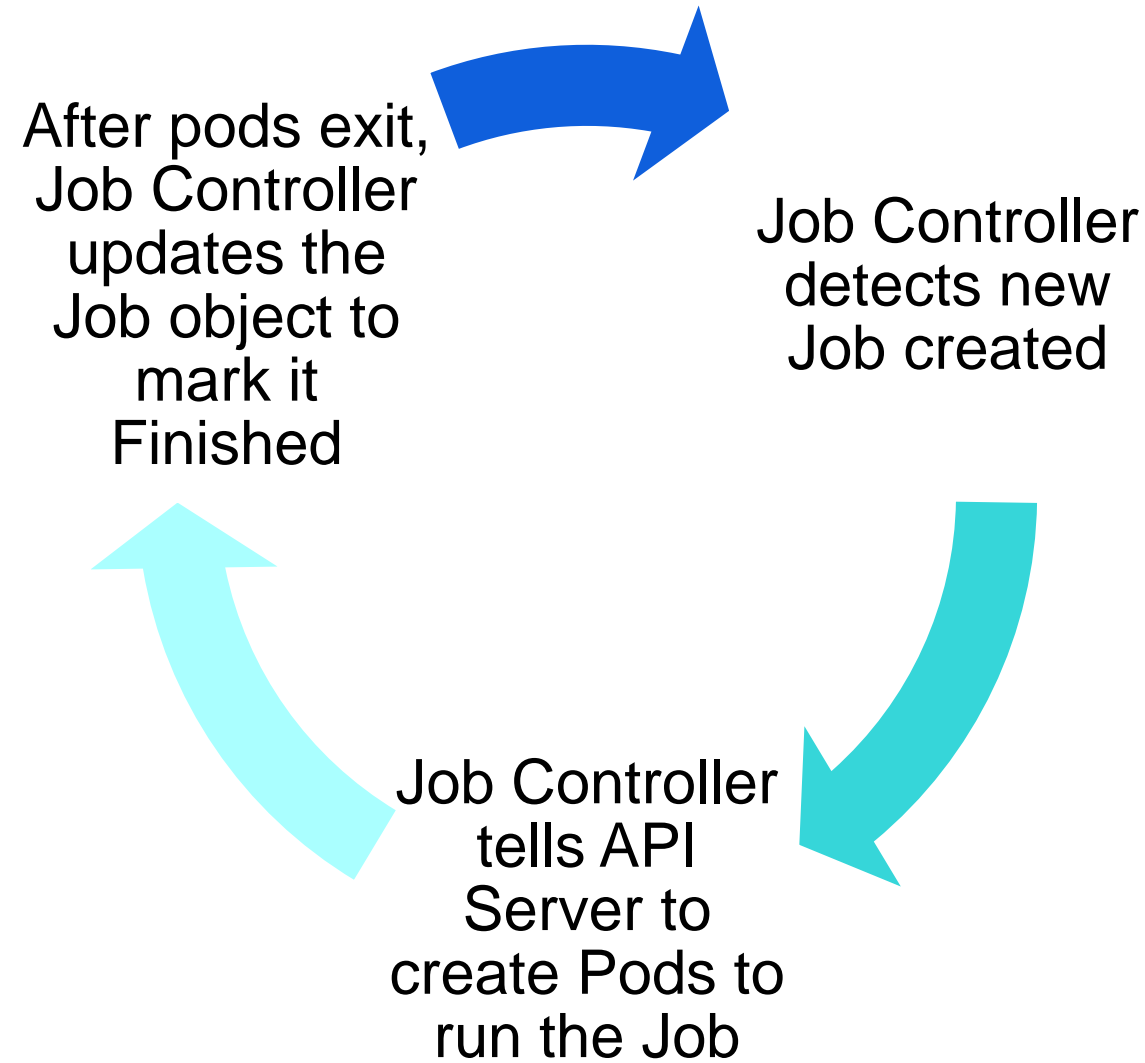
Types of workload resources:

- *Deployment*
- *ReplicaSet*
- *DaemonSet*
- *Job*
- *CronJob*
- *StatefulSet* (requires concepts about persistent storage and will be discussed in a later lesson).

When you submit a workload resource, the K8s control plane configures a corresponding *controller*.

Controllers run *control loops*: infinite loops in which the controller periodically checks if the status of the resource corresponds to the declared desired state, and if it doesn't match, it tries to correct the situation.

Control Loop (example of the Job Controller)



<https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
<https://kubernetes.io/docs/concepts/architecture/controller/>

Spec and status fields

Most Kubernetes objects have two fields: *spec* and *status*.

The *spec* is set when you create the object, and it describes the object **desired state**.

You create the *spec* when you write the YAML file that described the object.

The *status* describes the **current state** of the object.

When you get the object at runtime in YAML format, you see also the *status*.

The controller continually tries to make the object actual state match the desired state.

Deployment

A deployment is a type of workload that is suitable for running stateless applications, such as web servers.

You can create deployments with the CLI or with a YAML file.

In a Deployment YAML file, you describe the desired number of replicas and the specification of the containers that you want to run.

The Deployment Controller ensures that the cluster realizes the desired state, by creating or deleting pods as required to match the desired number of replicas specified by the Deployment.

Creating deployments from the command line

- To create a deployment from an image:

```
kubectl create deployment --image <image-name> <deployment-name>
```

```
lara@kube-master-gui:~$ kubectl create deployment --image nginx nginx-app
deployment.apps/nginx-app created
```

- To list the deployment:

```
kubectl get deployment <deployment-name>
```

```
lara@kube-master-gui:~$ kubectl get deployment nginx-app
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-app     0/1     1            0           27s
```

```
lara@kube-master-gui:~$ kubectl get deployment nginx-app
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-app     1/1     1            1           75s
```

- To list the pods (created as part of the deployment):

```
kubectl get pods
```

```
lara@kube-master-gui:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-app-d6ff45774-nfbl2          1/1     Running   0           6m6s
```

Describing deployments

You can describe the deployment using:

```
kubectl describe deployment <deployment-name>
```

As you can see there are no ports, hence it won't be possible to access the nginx server from a web browser. You can use `--port=80` to publish the port when creating the deployment.

```
lara@kube-master-gui:~$ kubectl describe deployment nginx-app
Name:                nginx-app
Namespace:           default
CreationTimestamp:    Sun, 03 Jan 2021 12:43:43 +0100
Labels:              app=nginx-app
Annotations:         deployment.kubernetes.io/revision: 1
Selector:            app=nginx-app
Replicas:            1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:        RollingUpdate
MinReadySeconds:     0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx-app
  Containers:
    nginx:
      Image:      nginx
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available      True    MinimumReplicasAvailable
    Progressing    True    NewReplicaSetAvailable
OldReplicaSets:  <none>
NewReplicaSet:   nginx-app-d6ff45774 (1/1 replicas created)
Events:
  Type           Reason             Age   From                  Message
  ----           -
  Normal        ScalingReplicaSet   53s   deployment-controller  Scaled up replica set nginx-app-d6ff45774 to 1
```

Using YAML descriptors in Kubernetes

- To see the descriptor of the deployment in yaml format:

```
kubectl get deployment <deployment-name> -o yaml
```

```
kubectl get deployment nginx-app -o yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  ....
  labels:
    app: nginx-app
  managedFields:
  ....
  - apiVersion: apps/v1
  ...
  name: nginx-app
  namespace: default
  ...
spec:
  progressDeadlineSeconds: 600
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      app: nginx-app
```

```
....
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: nginx-app
    spec:
      containers:
        - image: nginx
          imagePullPolicy: Always
          name: nginx
          resources: {}
          terminationMessagePath: /dev/termination-log
          terminationMessagePolicy: File
      dnsPolicy: ClusterFirst
      restartPolicy: Always
      schedulerName: default-scheduler
      securityContext: {}
      terminationGracePeriodSeconds: 30
  status:
  ...
```


YAML syntax 1/3

The YAML specification can be found here: <https://yaml.org/spec/1.2/spec.html>

Most important rules:

- Indentation denotes structure. A nested block is indented with respect to its parent and so on, hierarchically. Indentation must be made with spaces, you cannot use tabs. The number of spaces used to indent a block does not have any meaning but must be used consistently. A typical document indents sub-blocks by 2 spaces.
- Comments are introduced by the pound sign “#” and continue for the rest of the line. They can appear anywhere.
- The elements of a sequence of scalars (strings, numbers) are introduced by a minus sign and a space:
 - element1
 - element2
 - element3
- Maps are defined separating the key from the value with a colon followed by a space. This is to ensure that strings containing columns don't have to be quoted (e.g. `https://localhost`):
 - key1: value1
 - key2: value2
- Mapping scalar keys to sequence values combines the two constructs above:
 - key1:
 - value1
 - value2

YAML syntax 2/3

- Sequence of mappings:

-

```
key1: value1
key2: value2
```

-

```
key1: value3
key2: value4
```

- In compact form a sequence of mappings can be written as:

- key1: value1

```
key2: value2
```

- key1: value3

```
key2: value4
```

- Sequences can also be written in square brackets, with comma separated elements:

```
[element1, element2, element3]
```

- Mappings can also be written in curly braces, with comma separated `key: value` pairs:

```
{key1: value1, key2: value2}
```

- Strings may be quoted. Single quotes don't allow for escape sequences. Double quotes may contain special characters escaped with `"\"`, such as `"\n"`.
- Three dashes may be used to indicate the start of a document `---`
- Three dots may be used to indicate the end of a document `...`

YAML syntax 3/3

- YAML may use tags to define the types of scalars. Tags are introduced by an exclamation mark "!". However in Kubernetes you will mainly see untagged scalars. A YAML document can use the following data types:

```
my_integer: 10
my_float: 1.23015e+3
my_Boolean: true
my_multiline_string:
  This is the first line
  This is the second line
my_double_quoted_string: " # This is not a comment"
my_single_quoted_string: '"Nested quotes" can go inside single quotes'
my_date: 2021-01-01
```

- Long literals can be written with multiple lines that preserve newlines:

```
my_key: |
  my_value_line_1
  my_value_line_2
  my_value_line_3
```

- Long literals can be written with multiple lines that get folded (newline replaced by space)

```
my_key: >
  This is the beginning
  of a long sentence
  with just spaces between words
```

Creating a deployment using a yaml file

- A Kubernetes descriptor must have the fields:
 - `apiVersion`
 - `kind`
 - `metadata`
 - `spec`
 - The spec has the fields:
 - `selector`
 - `template`
- To create the deployment using the file, use:
`kubectl apply -f <file-name.yml>`
- To list deployments do:
`kubectl get deployments`
- To describe a deployment do:
`kubectl describe <deployment-name>`
- Describing the deployment does not provide the IP address of the pods, to see that you need to describe each individual pod.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx-app
  minReadySeconds: 5
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

Accessing nginx from a deployment, need to access specific pod ip

```
lara@kube-master-gui:~/k8s/deployments$ kubectl apply -f nginx-deployment.yml
deployment.apps/nginx-deployment created
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6f7d8d4d55-468x9   1/1     Running   0           11s
lara@kube-master-gui:~/k8s/deployments$ kubectl describe pod
Name:                               nginx-deployment-6f7d8d4d55-468x9
Namespace:                           default
Priority:                             0
Node:                               kube-master-gui/10.0.2.15
Start Time:                         Sun, 03 Jan 2021 16:49:04 +0100
Labels:                             app=nginx-app
                                     pod-template-hash=6f7d8d4d55
Annotations:                         cni.projectcalico.org/podIP: 10.1.96.179/32
                                     cni.projectcalico.org/podIPs: 10.1.96.179/32
Status:                             Running
IP:                                  10.1.96.179
IPs:
  IP:                                10.1.96.179
Controlled By:                       ReplicaSet/nginx-deployment-6f7d8d4d55
Containers:
  nginx:
    Container ID:                    containerd://bc4ea1e2d20c7a20d3bd21a9fd211eb019ae6c27e476f8bd0ba3c6819ce87e5c
    Image:                           nginx
    Image ID:                         docker.io/library/nginx@sha256:4cf620a5c81390ee209398ecc18e5fb9dd0f5155cd82adcbae532fec94006fb9
    Port:                             80/TCP
    Host Port:                        0/TCP
    State:                            Running
      Started:                        Sun, 03 Jan 2021 16:49:06 +0100
    Ready:                            True
```

```
lara@kube-master-gui:~/k8s/deployments$ curl -k http://10.1.96.179
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

ReplicaSets

There is a specific ReplicaSet object whose purpose is to maintain the desired number of instances of a pod defined by a template.

However, in general you don't need to create an explicit ReplicaSet object and you can simply specify the desired number of replicas in a Deployment object.

- *spec.replicas* defines the desired number of replicas
- *spec.selector.matchLabels* defines the logic for identifying which pods are part of this deployment
- *spec.metadata.labels* attaches the label *app: nginx-app* to every pod created from this deployment

```
lara@kube-master-gui:~/k8s/deployments$ kubectl apply -f nginx-3-deployment.yml
deployment.apps/nginx-deployment created
lara@kube-master-gui:~/k8s/deployments$ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    3/3     3            3           6s
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6f7d8d4d55-gxb7l   1/1     Running   0           12s
nginx-deployment-6f7d8d4d55-s699f   1/1     Running   0           12s
nginx-deployment-6f7d8d4d55-dggc7   1/1     Running   0           12s
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
labels:
  app: nginx-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx-app
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

Scaling the number of pods by editing a deployment

Kubectl edit allows you to interactively modify the yaml definition of an object.

```
kubectl edit deployment <deployment-name>
```

By default kubectl uses the vi editor on Linux. If you prefer the nano editor, set the following environment variable:

```
export KUBE_EDITOR=nano
```

When editing the file, change the number of replicas from 3 to 5.

When saving the edited file, nano proposes a file name that differs from the original file you used to create the deployment. This is fine, accept the default.

Then check how many pods belong to the deployment with a label selector that matches the label specified in the deployment:

```
kubectl get pods -l app=nginx-app
lara@k8s-master-gui:~/k8s/deployments$ kubectl edit deployment nginx-deployment
deployment.apps/nginx-deployment edited
lara@k8s-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-6f7d8d4d55-gxb7l	1/1	Running	0	5m44s
nginx-deployment-6f7d8d4d55-s699f	1/1	Running	0	5m44s
nginx-deployment-6f7d8d4d55-dggc7	1/1	Running	0	5m44s
nginx-deployment-6f7d8d4d55-ktff2	1/1	Running	0	8s
nginx-deployment-6f7d8d4d55-4dxdt	1/1	Running	0	8s

Scaling the number of pods from the command line

You can also scale the number of pods using an imperative command:

```
kubectl scale deployment <deployment-name> --replicas=5
```

```
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6f7d8d4d55-dlfh5   1/1     Running   0           3m18s
nginx-deployment-6f7d8d4d55-kvl2r   1/1     Running   0           3m18s
nginx-deployment-6f7d8d4d55-6b9tz   1/1     Running   0           3m18s
lara@kube-master-gui:~/k8s/deployments$ kubectl scale deployment.apps/nginx-deployment --replicas=5
deployment.apps/nginx-deployment scaled
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
NAME                                READY   STATUS             RESTARTS   AGE
nginx-deployment-6f7d8d4d55-dlfh5   1/1     Running            0           3m30s
nginx-deployment-6f7d8d4d55-kvl2r   1/1     Running            0           3m30s
nginx-deployment-6f7d8d4d55-6b9tz   1/1     Running            0           3m30s
nginx-deployment-6f7d8d4d55-9r9gf   0/1     ContainerCreating  0           3s
nginx-deployment-6f7d8d4d55-skr94   1/1     Running            0           3s
```

Scaling deployments based on conditions

You can scale deployments automatically based on certain conditions being verified. Example:

```
kubectl autoscale deployment <deployment-name> --min=7 --max=8 --cpu-percent=60
```

This command causes the creation of a HorizontalPodAutoscaler object (hpa).

If the target average CPU utilization (represented as a percent of requested CPU) over all the pods is exceeded, this autoscaler will try to increase the number of pods up to the maximum value.

You can create, delete, get, describe HorizontalPodAutoscaler (hpa) objects.

```
lara@kube-master-gui:~/k8s/deployments$ kubectl get hpa
No resources found in default namespace.
lara@kube-master-gui:~/k8s/deployments$ kubectl autoscale deployment nginx-deployment --min=7 --max=8 --cpu-percent=60
horizontalpodautoscaler.autoscaling/nginx-deployment autoscaled
lara@kube-master-gui:~/k8s/deployments$ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    5/7     7            5           38m
lara@kube-master-gui:~/k8s/deployments$ kubectl get hpa
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
nginx-deployment    Deployment/nginx-deployment  <unknown>/60%   7         8         7          19s
lara@kube-master-gui:~/k8s/deployments$ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    7/7     7            7           38m
```

Deleting deployments

To delete deployments, use the command:

```
kubectl delete deployment <deployment-name>
```

You can then verify that all the pods are automatically deleted.

This is an example of use of Garbage Collector: when the parent of an object no longer exists, the object itself gets deleted automatically.

It takes some times for all the pods to be deleted, and you will see them in Terminating state for some seconds.

```
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6f7d8d4d55-gxb7l   1/1     Running   0           5m44s
nginx-deployment-6f7d8d4d55-s699f   1/1     Running   0           5m44s
nginx-deployment-6f7d8d4d55-dggc7   1/1     Running   0           5m44s
nginx-deployment-6f7d8d4d55-ktff2   1/1     Running   0           8s
nginx-deployment-6f7d8d4d55-4xdxt   1/1     Running   0           8s
lara@kube-master-gui:~/k8s/deployments$ kubectl delete deployment nginx-deployment
deployment.apps "nginx-deployment" deleted
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6f7d8d4d55-gxb7l   1/1     Terminating   0           14m
nginx-deployment-6f7d8d4d55-s699f   1/1     Terminating   0           14m
nginx-deployment-6f7d8d4d55-dggc7   1/1     Terminating   0           14m
nginx-deployment-6f7d8d4d55-ktff2   1/1     Terminating   0           8m37s
nginx-deployment-6f7d8d4d55-4xdxt   1/1     Terminating   0           8m37s
lara@kube-master-gui:~/k8s/deployments$ kubectl get pods -l app=nginx-app
No resources found in default namespace.
```

DaemonSet

A DaemonSet ensures that pods run on each node of the cluster, even if nodes are added at a later time.

A typical example of use for a DaemonSet is a process that collects log on each Node of the Cluster.

Fluentd is such a tool that can be used for log collection. On microk8s you can enable the fluentd add-on.

Then you will see that there is a DaemonSet called fluentd-...

```
lara@kube-master-gui:~/k8s/deployments$ microk8s enable fluentd
Enabling Fluentd-Elasticsearch
Labeling nodes
node/kube-master-gui labeled
Addon dns is already enabled.
--allow-privileged=true
service/elasticsearch-logging created
serviceaccount/elasticsearch-logging created
clusterrole.rbac.authorization.k8s.io/elasticsearch-logging created
clusterrolebinding.rbac.authorization.k8s.io/elasticsearch-logging created
statefulset.apps/elasticsearch-logging created
configmap/fluentd-es-config-v0.2.0 created
serviceaccount/fluentd-es created
clusterrole.rbac.authorization.k8s.io/fluentd-es created
clusterrolebinding.rbac.authorization.k8s.io/fluentd-es created
daemonset.apps/fluentd-es-v3.0.2 created
deployment.apps/kibana-logging created
service/kibana-logging created
Fluentd-Elasticsearch is enabled
```

```
lara@kube-master-gui:~/k8s/deployments$ kubectl get daemonset --all-namespaces
```

NAMESPACE	NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-system	calico-node	1	1	1	1	1	kubernetes.io/os=linux	3d6h
kube-system	fluentd-es-v3.0.2	1	1	0	1	0	<none>	2m20s

Job

The purpose of the Job object is to ensure that a specific task completes even if the pod that should run the task fails. In the case of standalone pod, if the pod fails it will not be restarted, and so the task won't complete.

A Job solves this problem by ensuring that the task is completed, so if the node where the pod is running fails, the scheduler will launch the pod on another node so that the task can be completed.

Jobs can run pods sequentially or in parallel.

The metadata attribute **parallelism** indicates how many pods should run in parallel.

The metadata attribute **completions** indicates how many pods should run in total before the job is complete.

The metadata attribute **backoffLimit** indicates how many pods should fail before the job itself is failed permanently.

Job example

```
apiVersion: batch/v1
kind: Job
metadata:
  name: job1
spec:
  template:
    spec:
      containers:
      - name: job1container
        image: alpine
        imagePullPolicy: IfNotPresent
        command: ['ash', '-c', 'echo Job Pod is Running ;
sleep 10']
      restartPolicy: Never
backoffLimit: 4
completions: 3
parallelism: 2
```

```
lara@kube-master-gui:~/k8s/jobs$ kubectl apply -f job1.yml
job.batch/job1 created
lara@kube-master-gui:~/k8s/jobs$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
job1-6mnnl    1/1     Running   0           6s
job1-tl85g    1/1     Running   0           6s
lara@kube-master-gui:~/k8s/jobs$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
job1-6mnnl    1/1     Running   0          11s
job1-tl85g    1/1     Running   0          11s
lara@kube-master-gui:~/k8s/jobs$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
job1-6mnnl    0/1     Completed  0          24s
job1-tl85g    0/1     Completed  0          24s
job1-tkp7q    1/1     Running   0          12s
lara@kube-master-gui:~/k8s/jobs$ kubectl get jobs
NAME    COMPLETIONS  DURATION  AGE
job1    3/3          25s       37s
lara@kube-master-gui:~/k8s/jobs$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
job1-6mnnl    0/1     Completed  0          41s
job1-tl85g    0/1     Completed  0          41s
job1-tkp7q    0/1     Completed  0          29s
```

You can also more specifically execute: `kubectl get pods -l job-name=job1`

CronJob

A CronJob causes a pod to execute at specific times determined by the `schedule` parameter that follows the syntax of Unix cron:
<https://en.wikipedia.org/wiki/Cron>

The cronjob in the example executes every minute.

If you look for the pods you can see that a new pod completes every minute.

Note that CronJob became stable in k8s 1.21.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: cronjob1
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: cronjob1container
              image: alpine
              imagePullPolicy: IfNotPresent
              command: ['ash', '-c', 'echo Job Pod is Running ; sleep 10']
              restartPolicy: OnFailure
```

```
lara@k8s-master-gui:~/k8s/jobs$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cronjob1-1609713180-qlbmc	0/1	Completed	0	2m11s
cronjob1-1609713240-h2kw2	0/1	Completed	0	71s
cronjob1-1609713300-m82jl	1/1	Running	0	9s

Garbage Collector

- The Garbage Collector is responsible for deleting objects when their parent no longer exists.
- This is implemented by setting the field `ownerReferences` in the child object.
- In this example, a Pod created as part of a Job has `ownerReferences` set to the parent Job.
- **Cascading deletion** causes the child to be deleted when the parent is deleted, and it can happen in **foreground** or **background**.
- In **foreground**, the parent is marked for deletion. Then all objects with `blockOwnerDeletion: true` are deleted and finally, the parent is deleted.
- In **background**, first the parent is deleted and then all the children are found and deleted.

```
lara@kube-master-gui:~/k8s/jobs$ kubectl get pod job1-6mnn1 -o yaml
apiVersion: v1
kind: Pod
metadata:
  ....
  name: job1-6mnn1
  namespace: default
  ownerReferences:
  - apiVersion: batch/v1
    blockOwnerDeletion: true
    controller: true
    kind: Job
    name: job1
    uid: d5de87ae-3f50-445a-bfbd-4b5cb956f2c4
  resourceVersion: "263565"
  selfLink: /api/v1/namespaces/default/pods/job1-6mnn1
  uid: 14d9d3ed-192d-46cc-a0f5-7bd4dc2e6cc4
spec:
  containers:
  ...
```


HCLSoftware

hcltechsw.com